

Experiment No.4

Title : Use Autoencoder to implement anomaly detection

In [19]:

```
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras import Model, Sequential
from tensorflow.keras.layers import Dense, Dropout
from sklearn.model_selection import train_test_split
from tensorflow.keras.losses import MeanSquaredLogarithmicError
```

Load the data

In [2]:

```
# Download the dataset
PATH_TO_DATA = 'http://storage.googleapis.com/download.tensorflow.org/data/ecg.csv'
data = pd.read_csv(PATH_TO_DATA, header=None)
print(data.shape)
data.head()
```

(4998, 141)

Out[2]:

	0	1	2	3	4	5	6	7	
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376041	-3.474986	-2.181408	-1.818286	-1.2501
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.992258	-0.7541
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187449	-3.151462	-1.742940	-1.490659	-1.1831
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.671131	-1.3331
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423	-1.5941

5 rows × 141 columns

Split the data for training and testing

In [3]:

```
# Last column is the target
# 0 = anomaly, 1 = normal
TARGET = 140

features = data.drop(TARGET, axis=1)
target = data[TARGET]

x_train, x_test, y_train, y_test = train_test_split(
    features, target, test_size=0.2, stratify=target
)
```

Scale the data using MinMaxScaler

In [4]:

```
# use case is novelty detection so use only the normal data
# for training
train_index = y_train[y_train == 1].index
train_data = x_train.loc[train_index]
```

In [20]:

```
min_max_scaler = MinMaxScaler(feature_range=(0, 1))
x_train_scaled = min_max_scaler.fit_transform(train_data.copy())
x_test_scaled = min_max_scaler.transform(x_test.copy())
```

Build an AutoEncoder model

In [6]:

```

# create a model by subclassing Model class in tensorflow
class AutoEncoder(Model):
    """
    Parameters
    -----
    output_units: int
        Number of output units

    code_size: int
        Number of units in bottle neck
    """

    def __init__(self, output_units, code_size=8):
        super().__init__()
        self.encoder = Sequential([
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(code_size, activation='relu')
        ])
        self.decoder = Sequential([
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(output_units, activation='sigmoid')
        ])

    def call(self, inputs):
        encoded = self.encoder(inputs)
        decoded = self.decoder(encoded)
        return decoded

```

In [7]:

```

model = AutoEncoder(output_units=x_train_scaled.shape[1])
# configurations of model
model.compile(loss='msle', metrics=['mse'], optimizer='adam')

history = model.fit(
    x_train_scaled,
    x_train_scaled,
    epochs=20,
    batch_size=512,
    validation_data=(x_test_scaled, x_test_scaled)
)

```

Epoch 1/20

5/5 ————— 2s 38ms/step - loss: 0.0107 - mse: 0.0241 - val_loss: 0.0129 - val_mse: 0.0298

Epoch 2/20

5/5 ————— 0s 8ms/step - loss: 0.0097 - mse: 0.0218 - val_loss: 0.0127 - val_mse: 0.0292

Epoch 3/20

5/5 ————— 0s 8ms/step - loss: 0.0085 - mse: 0.0192 - val_loss: 0.0121 - val_mse: 0.0279

Epoch 4/20

5/5 ————— 0s 8ms/step - loss: 0.0073 - mse: 0.0165 - val_loss: 0.0116 - val_mse: 0.0267

Epoch 5/20

5/5 ————— 0s 9ms/step - loss: 0.0062 - mse: 0.0140 - val_loss: 0.0107 - val_mse: 0.0246

Epoch 6/20

5/5 ————— 0s 8ms/step - loss: 0.0053 - mse: 0.0119 - val_loss: 0.0098 - val_mse: 0.0226

Epoch 7/20

5/5 ————— 0s 8ms/step - loss: 0.0047 - mse: 0.0106 - val_loss: 0.0094 - val_mse: 0.0217

Epoch 8/20

5/5 ————— 0s 9ms/step - loss: 0.0044 - mse: 0.0100 - val_loss: 0.0092 - val_mse: 0.0214

Epoch 9/20

5/5 ————— 0s 9ms/step - loss: 0.0042 - mse: 0.0093 - val_loss: 0.0091 - val_mse: 0.0212

Epoch 10/20

5/5 ————— 0s 9ms/step - loss: 0.0042 - mse: 0.0094 - val_loss: 0.0091 - val_mse: 0.0211

Epoch 11/20

5/5 ————— 0s 8ms/step - loss: 0.0040 - mse: 0.0090 - val_loss: 0.0091 - val_mse: 0.0210

Epoch 12/20

5/5 ————— 0s 8ms/step - loss: 0.0039 - mse: 0.0087 - val_loss: 0.0090 - val_mse: 0.0209

Epoch 13/20

5/5 ————— 0s 9ms/step - loss: 0.0040 - mse: 0.0088 - val_loss: 0.0090 - val_mse: 0.0208

Epoch 14/20

5/5 ————— 0s 8ms/step - loss: 0.0038 - mse: 0.0085 - val_loss: 0.0089 - val_mse: 0.0207

Epoch 15/20

5/5 ————— 0s 11ms/step - loss: 0.0038 - mse: 0.0085 - val_loss: 0.0089 - val_mse: 0.0206

Epoch 16/20

5/5 ————— 0s 9ms/step - loss: 0.0037 - mse: 0.0084 - val_loss:

s: 0.0088 - val_mse: 0.0204

Epoch 17/20

5/5 ————— 0s 8ms/step - loss: 0.0036 - mse: 0.0081 - val_loss

s: 0.0087 - val_mse: 0.0202

Epoch 18/20

5/5 ————— 0s 9ms/step - loss: 0.0035 - mse: 0.0079 - val_loss

s: 0.0086 - val_mse: 0.0200

Epoch 19/20

5/5 ————— 0s 11ms/step - loss: 0.0033 - mse: 0.0075 - val_loss

s: 0.0085 - val_mse: 0.0197

Epoch 20/20


5/5 ————— 0s 9ms/step - loss: 0.0033 - mse: 0.0073 - val_loss

s: 0.0084 - val_mse: 0.0194

Plot history

In [8]:

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('MSLE Loss')
plt.legend(['loss', 'val_loss'])
plt.show()
```



Find threshold

In [9]:

```

def find_threshold(model, x_train_scaled):
    reconstructions = model.predict(x_train_scaled)
    # provides losses of individual instances
    reconstruction_errors = tf.keras.losses.msle(reconstructions, x_train_scaled)

    # threshold for anomaly scores
    threshold = np.mean(reconstruction_errors.numpy()) \
        + np.std(reconstruction_errors.numpy())
    return threshold

def find_threshold_method_two(model, x_train_scaled):
    # another method to find threshold
    reconstructions = model.predict(x_train_scaled)
    # provides losses of individual instances
    reconstruction_errors = tf.keras.losses.msle(reconstructions, x_train_scaled)

    threshold_2 = np.percentile(reconstruction_errors, 95)
    return threshold_2

def get_predictions(model, x_test_scaled, threshold):
    predictions = model.predict(x_test_scaled)
    # provides losses of individual instances
    errors = tf.keras.losses.msle(predictions, x_test_scaled)
    # 0 = anomaly, 1 = normal
    anomaly_mask = pd.Series(errors) > threshold
    preds = anomaly_mask.map(lambda x: 0.0 if x == True else 1.0)
    return preds

```

In [10]:

```

threshold = find_threshold(model, x_train_scaled)
print(f"Threshold method one: {threshold}")

threshold_2 = find_threshold_method_two(model, x_train_scaled)
print(f"Threshold method two: {threshold_2}")

```

```

73/73 ————— 0s 2ms/step
Threshold method one: 0.007880838475657386
73/73 ————— 0s 653us/step
Threshold method two: 0.011669582318263243

```

In [11]:

```

preds = get_predictions(model, x_test_scaled, threshold)
accuracy_score(preds, y_test)

```

```

32/32 ————— 0s 678us/step

```

Out[11]:

0.953

Tuning AutoEncoder using keras tuner

In [12]:

```
!pip install -U keras-tuner
```

Collecting keras-tuner

Obtaining dependency information for keras-tuner from https://files.pythonhosted.org/packages/db/5d/945296512980b0827e93418514c8be9236baa6f0a1e8ca8be3a2026665b0/keras_tuner-1.4.7-py3-none-any.whl.metadata (https://files.pythonhosted.org/packages/db/5d/945296512980b0827e93418514c8be9236baa6f0a1e8ca8be3a2026665b0/keras_tuner-1.4.7-py3-none-any.whl.metadata)

Downloading keras_tuner-1.4.7-py3-none-any.whl.metadata (5.4 kB)

Requirement already satisfied: keras in c:\users\admin\anaconda3\lib\site-packages (from keras-tuner) (3.4.1)

Requirement already satisfied: packaging in c:\users\admin\anaconda3\lib\site-packages (from keras-tuner) (23.1)

Requirement already satisfied: requests in c:\users\admin\anaconda3\lib\site-packages (from keras-tuner) (2.31.0)

Collecting kt-legacy (from keras-tuner)

Obtaining dependency information for kt-legacy from https://files.pythonhosted.org/packages/16/53/aca9f36da2516db008017db85a1f3cafaee0efc5fc7a25d94c909651792f/kt_legacy-1.0.5-py3-none-any.whl.metadata (https://files.pythonhosted.org/packages/16/53/aca9f36da2516db008017db85a1f3cafaee0efc5fc7a25d94c909651792f/kt_legacy-1.0.5-py3-none-any.whl.metadata)

Downloading kt_legacy-1.0.5-py3-none-any.whl.metadata (221 bytes)

Requirement already satisfied: absl-py in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (2.1.0)

Requirement already satisfied: numpy in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (1.24.3)

Requirement already satisfied: rich in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (13.7.1)

Requirement already satisfied: namex in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (0.0.8)

Requirement already satisfied: h5py in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (3.11.0)

Requirement already satisfied: optree in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (0.12.1)

Requirement already satisfied: ml-dtypes in c:\users\admin\anaconda3\lib\site-packages (from keras->keras-tuner) (0.4.0)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\admin\anaconda3\lib\site-packages (from requests->keras-tuner) (2.0.4)

Requirement already satisfied: idna<4,>=2.5 in c:\users\admin\anaconda3\lib\site-packages (from requests->keras-tuner) (3.4)

Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\admin\anaconda3\lib\site-packages (from requests->keras-tuner) (1.26.16)

Requirement already satisfied: certifi>=2017.4.17 in c:\users\admin\anaconda3\lib\site-packages (from requests->keras-tuner) (2023.7.22)

Requirement already satisfied: typing-extensions>=4.5.0 in c:\users\admin\anaconda3\lib\site-packages (from optree->keras->keras-tuner) (4.12.2)

Requirement already satisfied: markdown-it-py>=2.2.0 in c:\users\admin\anaconda3\lib\site-packages (from rich->keras->keras-tuner) (2.2.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\admin\anaconda3\lib\site-packages (from rich->keras->keras-tuner) (2.15.1)

Requirement already satisfied: mdurl~=0.1 in c:\users\admin\anaconda3\lib\site-packages (from markdown-it-py>=2.2.0->rich->keras->keras-tuner) (0.1.0)

Downloading keras_tuner-1.4.7-py3-none-any.whl (129 kB)

----- 0.0/129.1 kB ? eta -:--:--

--- 10.2/129.1 kB ? eta -:--:--

----- 129.1/129.1 kB 1.9 MB/s eta 0:

00:00

Downloading kt_legacy-1.0.5-py3-none-any.whl (9.6 kB)

Installing collected packages: kt-legacy, keras-tuner
 Successfully installed keras-tuner-1.4.7 kt-legacy-1.0.5

In [13]:

```
import kerastuner as kt

class AutoEncoderTuner(Model):

    def __init__(self, hp, output_units, code_size=8):
        super().__init__()
        dense_1_units = hp.Int('dense_1_units', min_value=16, max_value=72, step=4)
        dense_2_units = hp.Int('dense_2_units', min_value=16, max_value=72, step=4)
        dense_3_units = hp.Int('dense_3_units', min_value=16, max_value=72, step=4)
        dense_4_units = hp.Int('dense_4_units', min_value=16, max_value=72, step=4)
        dense_5_units = hp.Int('dense_5_units', min_value=16, max_value=72, step=4)
        dense_6_units = hp.Int('dense_6_units', min_value=16, max_value=72, step=4)

        self.encoder = Sequential([
            Dense(dense_1_units, activation='relu'),
            Dropout(0.1),
            Dense(dense_2_units, activation='relu'),
            Dropout(0.1),
            Dense(dense_3_units, activation='relu'),
            Dropout(0.1),
            Dense(code_size, activation='relu')
        ])
        self.decoder = Sequential([
            Dense(dense_4_units, activation='relu'),
            Dropout(0.1),
            Dense(dense_5_units, activation='relu'),
            Dropout(0.1),
            Dense(dense_6_units, activation='relu'),
            Dropout(0.1),
            Dense(output_units, activation='sigmoid')
        ])

    def call(self, inputs):
        encoded = self.encoder(inputs)
        decoded = self.decoder(encoded)
        return decoded

def build_model(hp):
    model = AutoEncoderTuner(hp, 140)
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])
    model.compile(
        loss='mse',
        optimizer=Adam(learning_rate=hp_learning_rate),
    )
    return model
```

C:\Users\ADMIN\AppData\Local\Temp\ipykernel_8352\3309659638.py:1: Deprecatio
 nWarning: `import kerastuner` is deprecated, please use `import keras_tuner`

`.

```
import kerastuner as kt
```


In [14]:

```
tuner = kt.Hyperband(
    build_model,
    objective='val_loss',
    max_epochs=20,
    factor=3,
    directory='autoencoder',
    project_name='tuning_autoencoder6'
)

tuner.search(
    x_train_scaled,
    x_train_scaled,
    epochs=20,
    batch_size=512,
    validation_data=(x_test_scaled, x_test_scaled)
)
```

Trial 30 Complete [00h 00m 04s]
val_loss: 0.008870153687894344

Best val_loss So Far: 0.007670378312468529
Total elapsedtime: 00h 01m 17s

In [15]:

```
hparams = [f'dense_{i}_units' for i in range(1,7)] + ['learning_rate']
best_hyperparams = tuner.get_best_hyperparameters()
for hps in hparams:
    print(f"{hps}: {best_hyperparams[0][hps]}")
```

dense_1_units: 64
dense_2_units: 48
dense_3_units: 28
dense_4_units: 48
dense_5_units: 48
dense_6_units: 68
learning_rate: 0.01

In [16]:

```
best_model = tuner.get_best_models()[0]
best_model.compile(loss='mse', optimizer=Adam(0.001))

best_model.fit(
    x_train_scaled,
    x_train_scaled,
    epochs=20,
    batch_size=512,
    validation_data=(x_test_scaled, x_test_scaled)
)
```

```
Epoch 1/20
5/5 ————— 1s 32ms/step - loss: 0.0022 - val_loss: 0.0076
Epoch 2/20
5/5 ————— 0s 8ms/step - loss: 0.0022 - val_loss: 0.0076
Epoch 3/20
5/5 ————— 0s 9ms/step - loss: 0.0021 - val_loss: 0.0075
Epoch 4/20
5/5 ————— 0s 8ms/step - loss: 0.0021 - val_loss: 0.0075
Epoch 5/20
5/5 ————— 0s 8ms/step - loss: 0.0021 - val_loss: 0.0074
Epoch 6/20
5/5 ————— 0s 8ms/step - loss: 0.0020 - val_loss: 0.0073
Epoch 7/20
5/5 ————— 0s 8ms/step - loss: 0.0020 - val_loss: 0.0072
Epoch 8/20
5/5 ————— 0s 10ms/step - loss: 0.0020 - val_loss: 0.0071
Epoch 9/20
5/5 ————— 0s 10ms/step - loss: 0.0020 - val_loss: 0.0071
Epoch 10/20
5/5 ————— 0s 11ms/step - loss: 0.0019 - val_loss: 0.0070
Epoch 11/20
5/5 ————— 0s 7ms/step - loss: 0.0019 - val_loss: 0.0069
Epoch 12/20
5/5 ————— 0s 8ms/step - loss: 0.0018 - val_loss: 0.0069
Epoch 13/20
5/5 ————— 0s 7ms/step - loss: 0.0018 - val_loss: 0.0068
Epoch 14/20
5/5 ————— 0s 7ms/step - loss: 0.0018 - val_loss: 0.0067
Epoch 15/20
5/5 ————— 0s 11ms/step - loss: 0.0018 - val_loss: 0.0066
Epoch 16/20
5/5 ————— 0s 7ms/step - loss: 0.0018 - val_loss: 0.0066
Epoch 17/20
5/5 ————— 0s 7ms/step - loss: 0.0018 - val_loss: 0.0066
Epoch 18/20
5/5 ————— 0s 7ms/step - loss: 0.0017 - val_loss: 0.0065
Epoch 19/20
5/5 ————— 0s 7ms/step - loss: 0.0017 - val_loss: 0.0065
Epoch 20/20
5/5 ————— 0s 7ms/step - loss: 0.0017 - val_loss: 0.0064
```

Out[16]:

```
<keras.src.callbacks.history.History at 0x25a4eb62290>
```

In [17]:

```
threshold_ = find_threshold(best_model, x_train_scaled)
preds_ = get_predictions(best_model, x_test_scaled, threshold_)
accuracy_score(preds_, y_test)
```

73/73  0s 1ms/step
32/32  0s 710us/step

Out[17]:

0.968